

Using Storm

v0.8.2



storm-project.net

The Idea

Hadoop for real-time stream processing.

Abstractions

- Tuples and Tuple Trees
 - Streams and Stream Groupings
 - Shuffle Grouping
 - Field Grouping
 - All Grouping
 - Spouts
 - Bolts
 - Topologies
-

Cluster Anatomy

- Nimbus
- Zookeeper Cluster
- Worker Nodes
 - Supervisor
 - Worker Process
 - Executors
 - Tasks

Features

- Admin Interface
 - Tuple dependencies (Tuple Tree)
 - Transactions (Transactional Topologies)
 - Multi-tenancy (Isolation scheduler)
 - Clojure DSL (and other languages)
-

Limitations

- No topology swapping (currently)
 - No automatic load balancing
 - Nimbus and Supervisors do not automatically failover
 - Work reassignment is done on timeout
 - Work replay is done from the Spout
-

Similar Tools

- Yahoo S4
 - Flume / FlumeBase
 - Riemann
-